Kafka

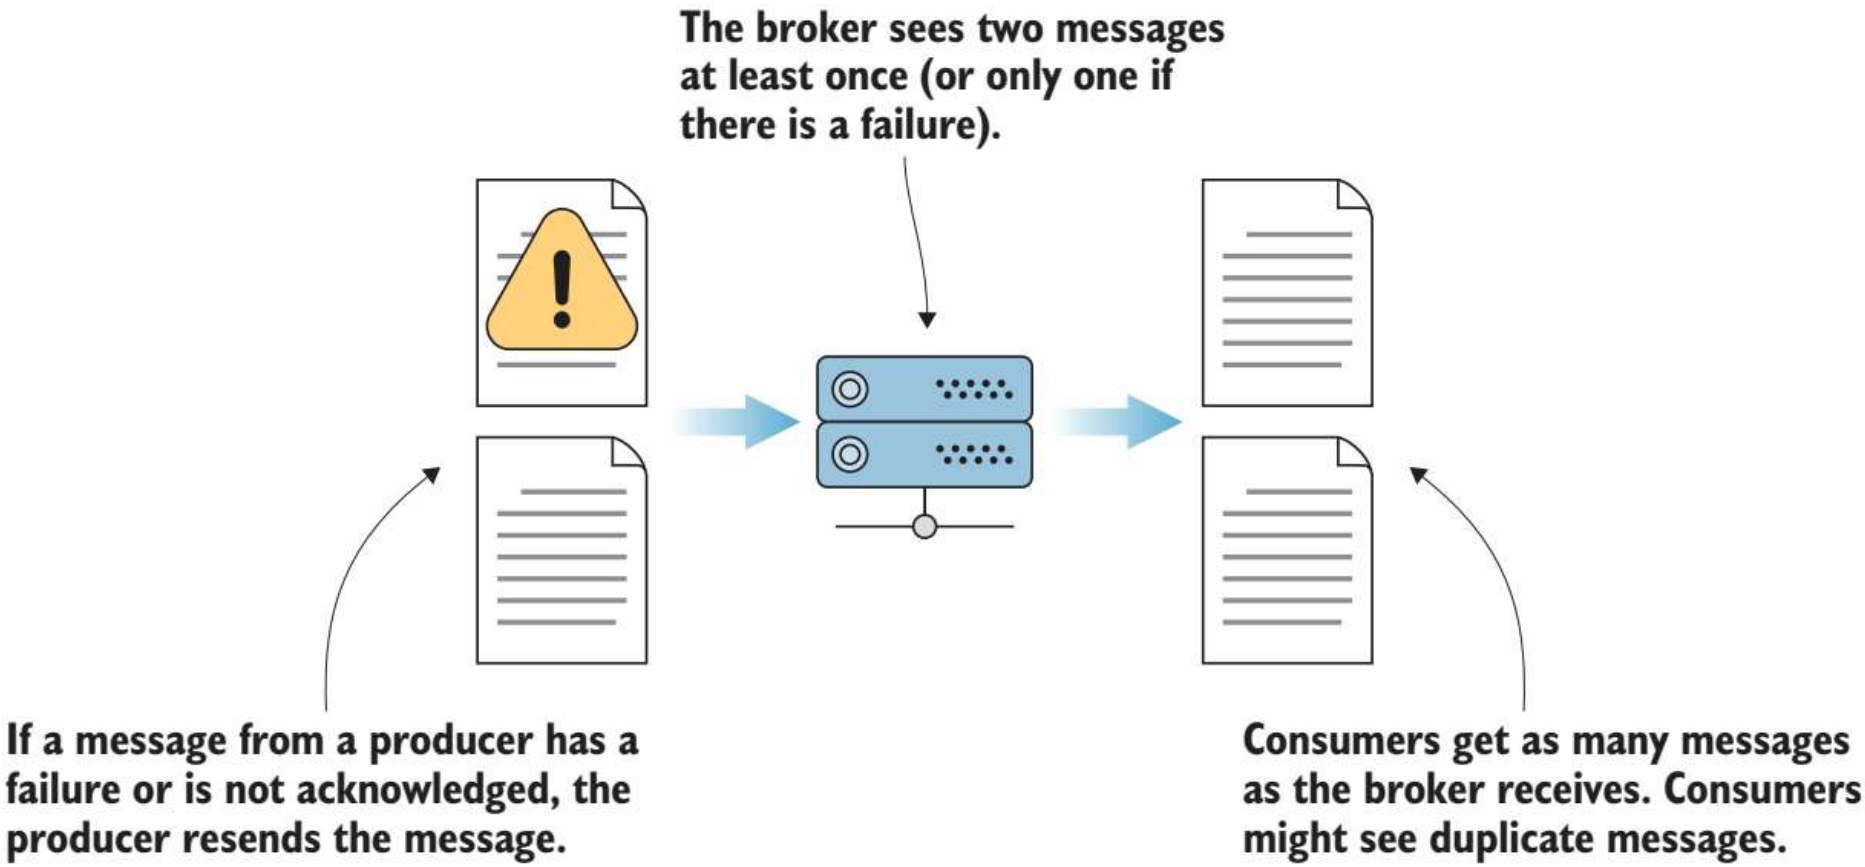

At-least-once message flow

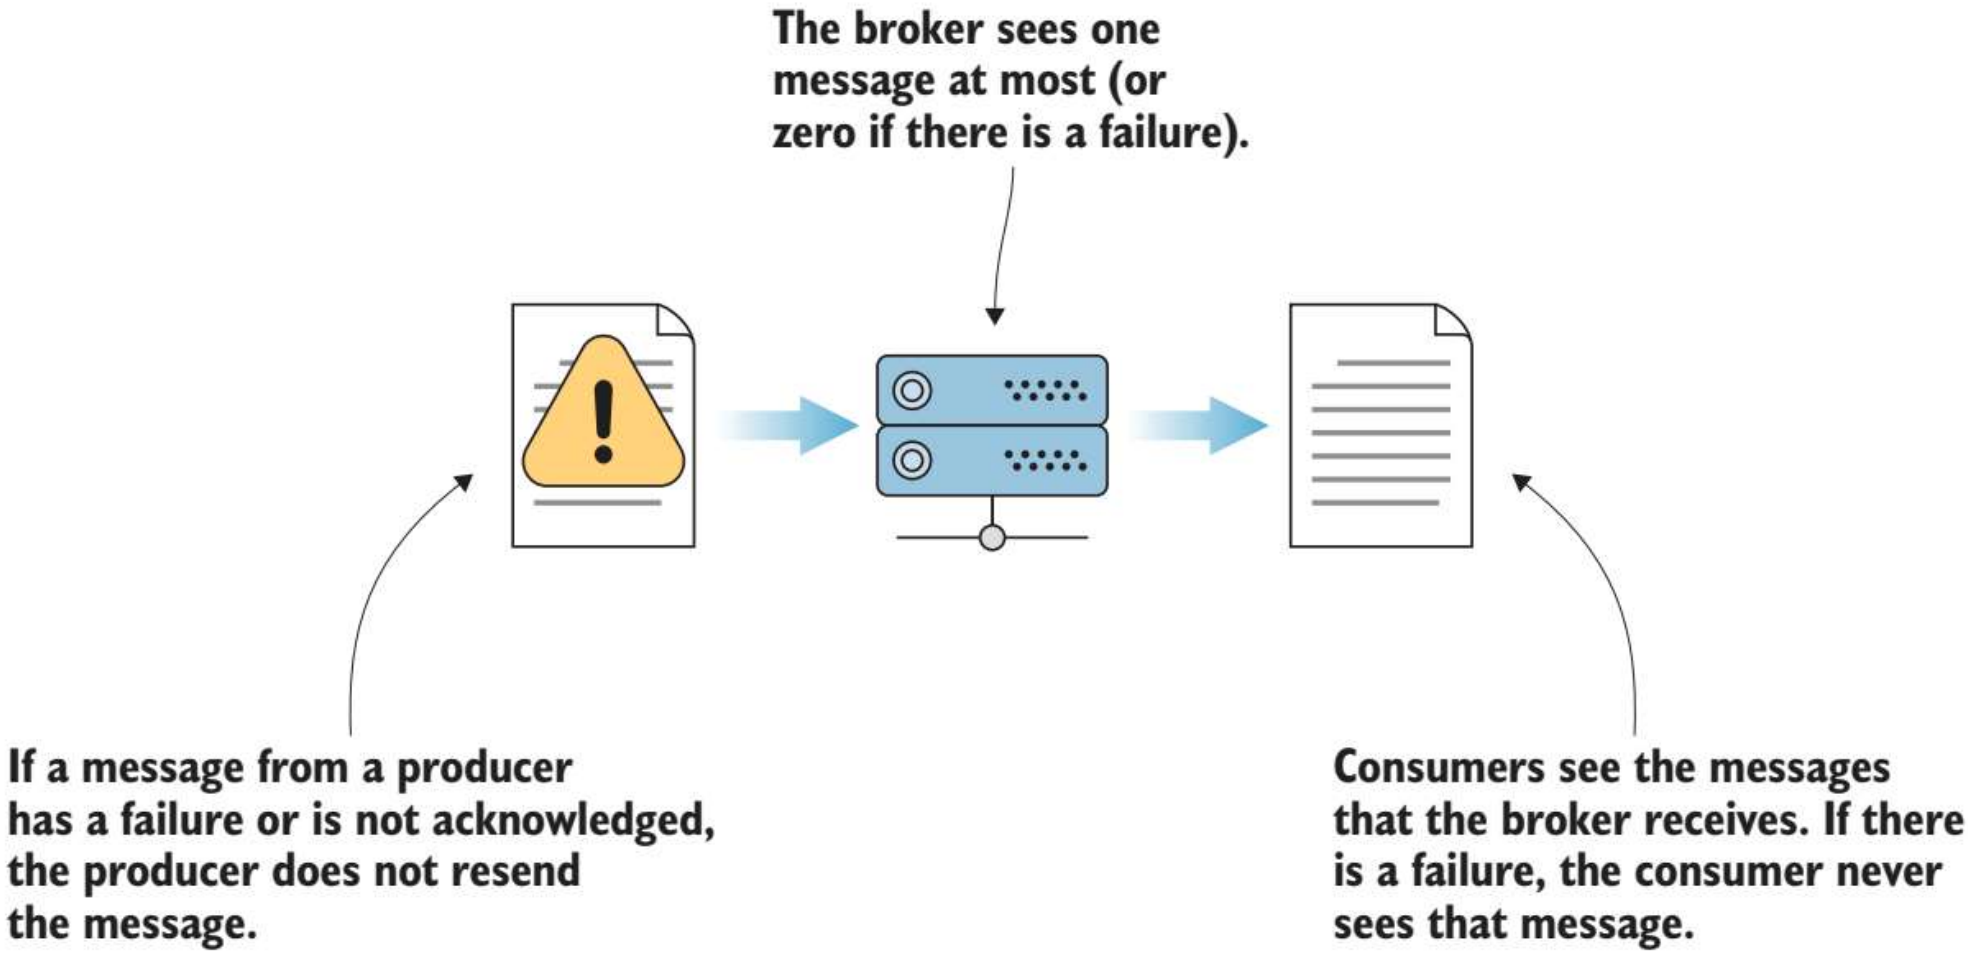

At-most-once message flow

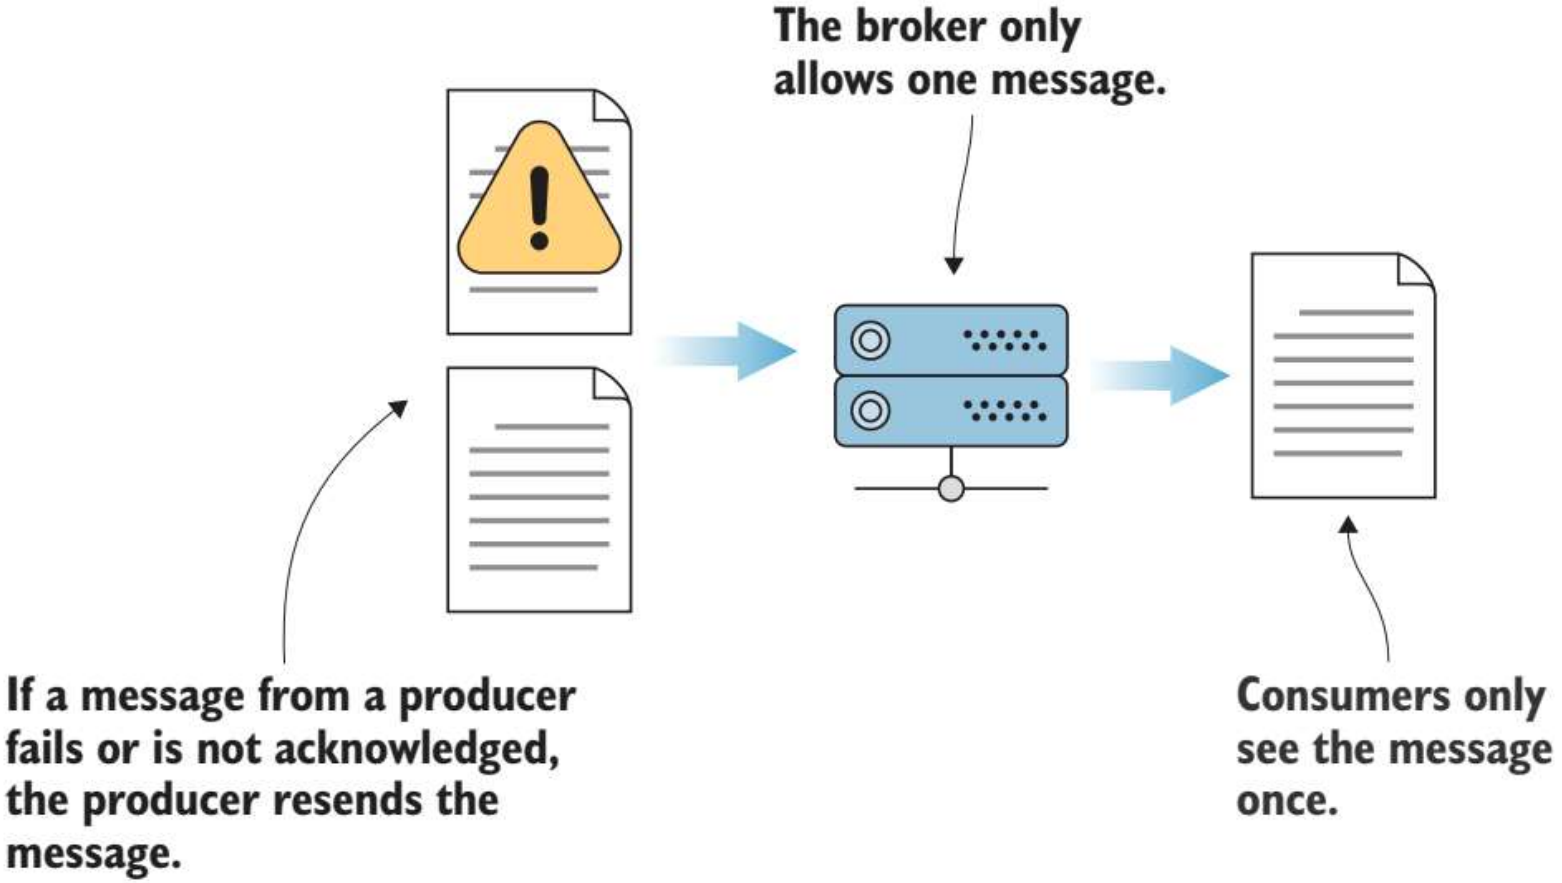

Exactly-once message flow

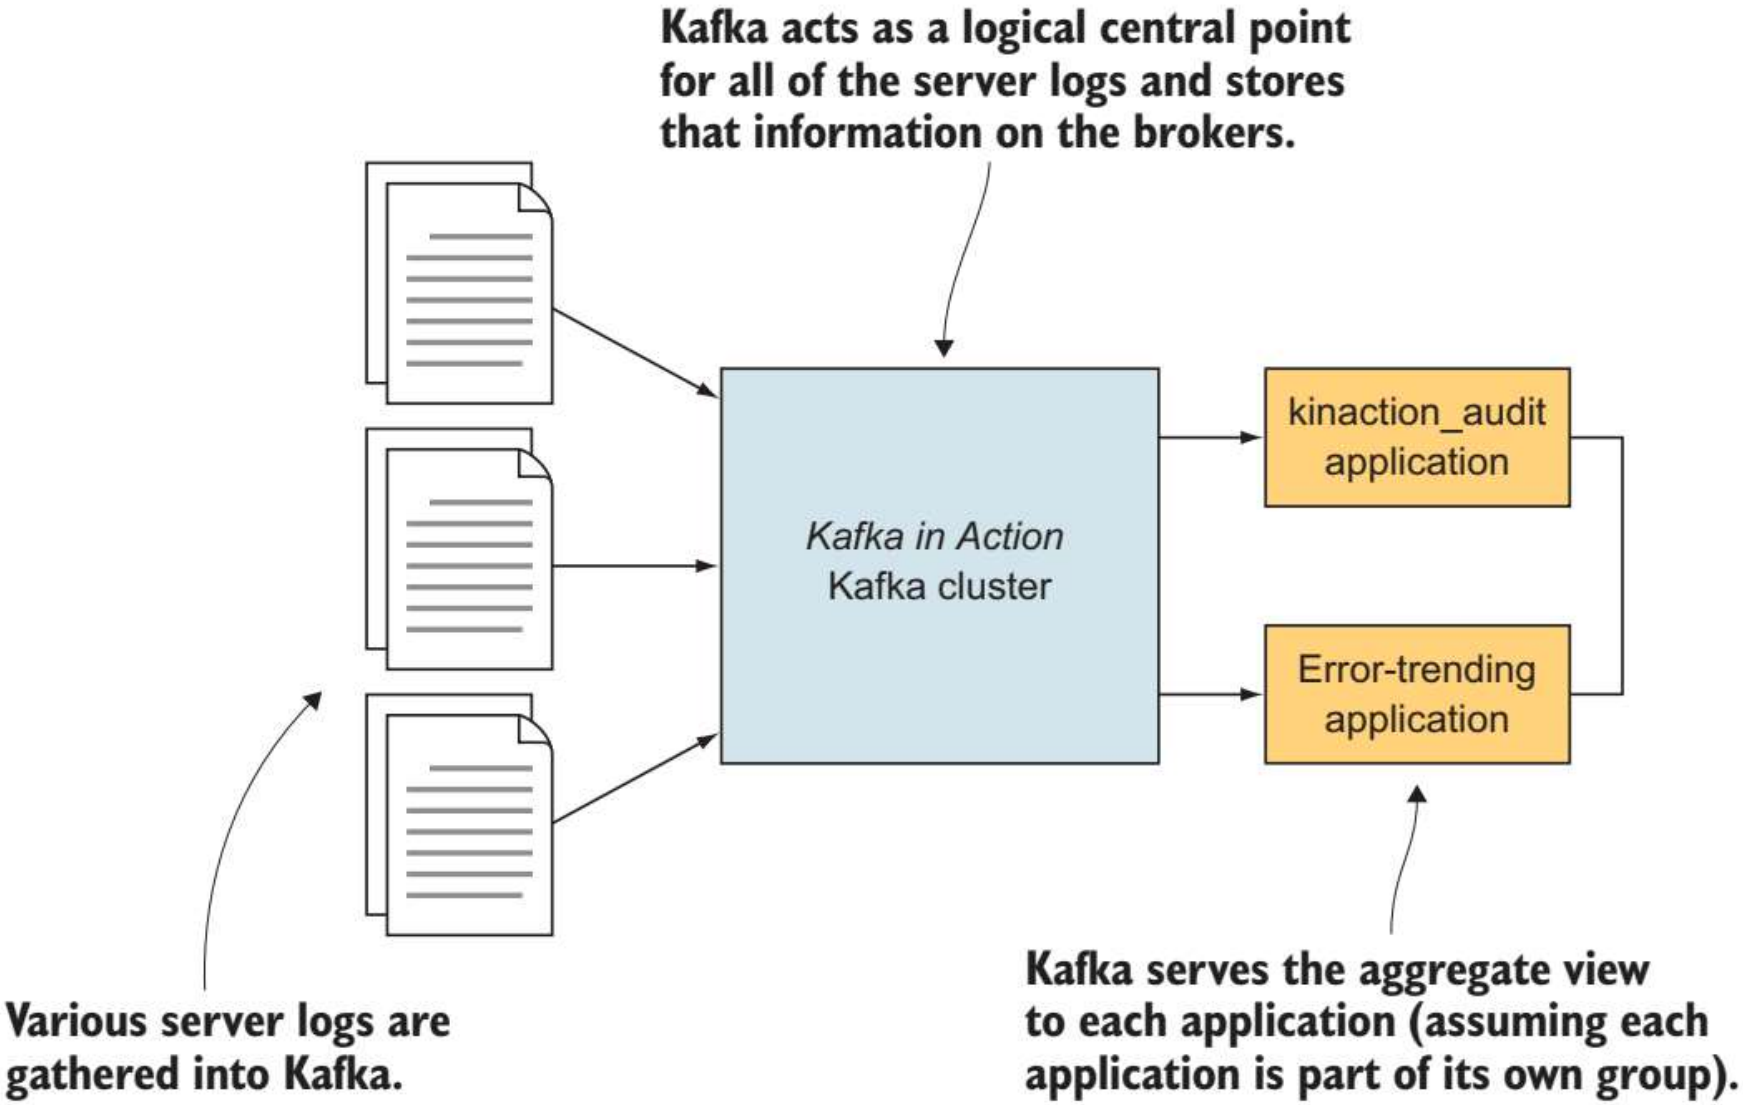

Kafka log aggregation

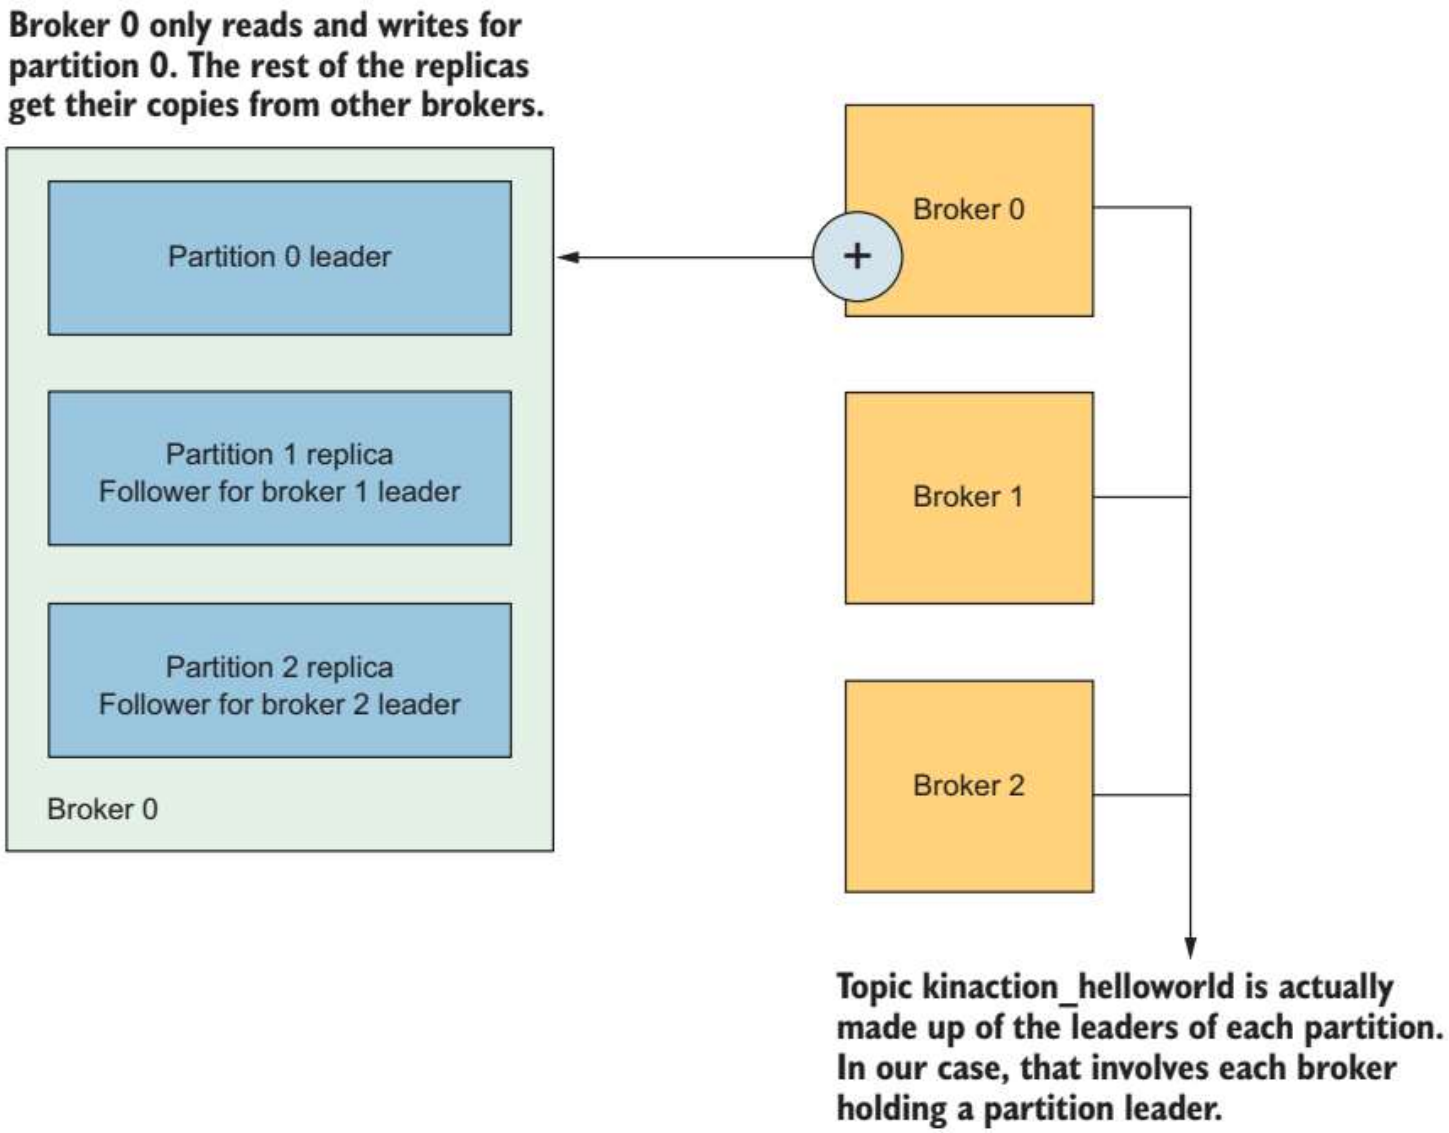

View of one broker

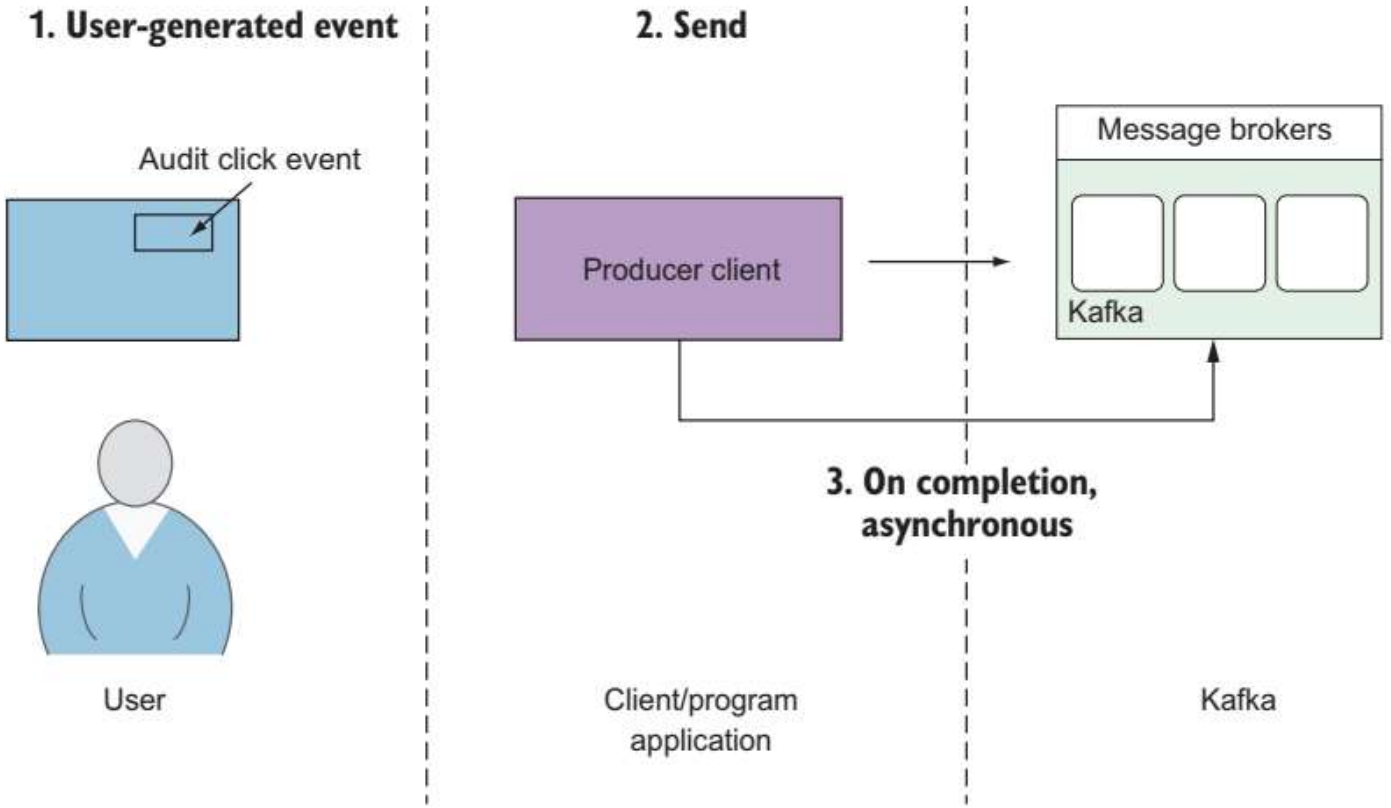

Producer example for user event

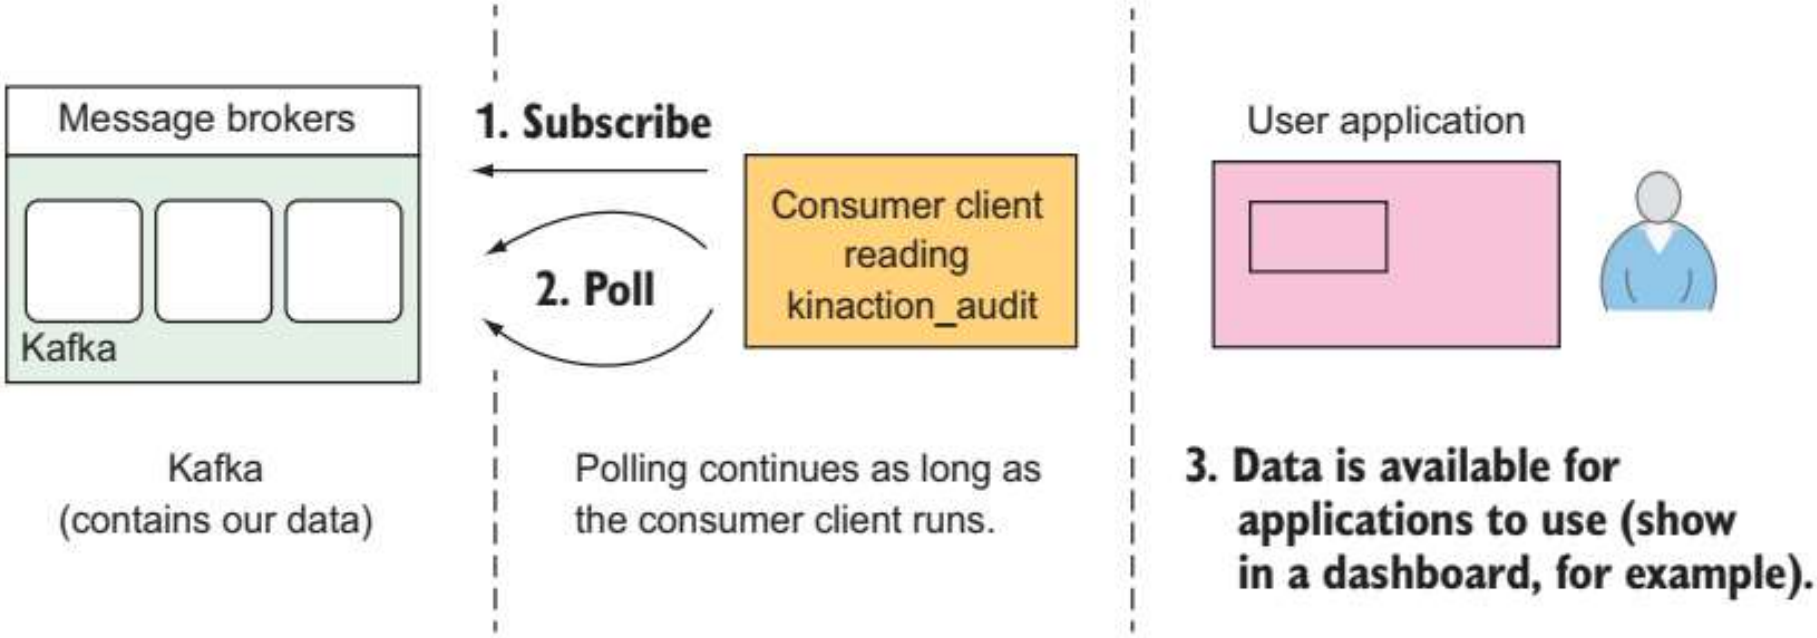

Consumer example flow

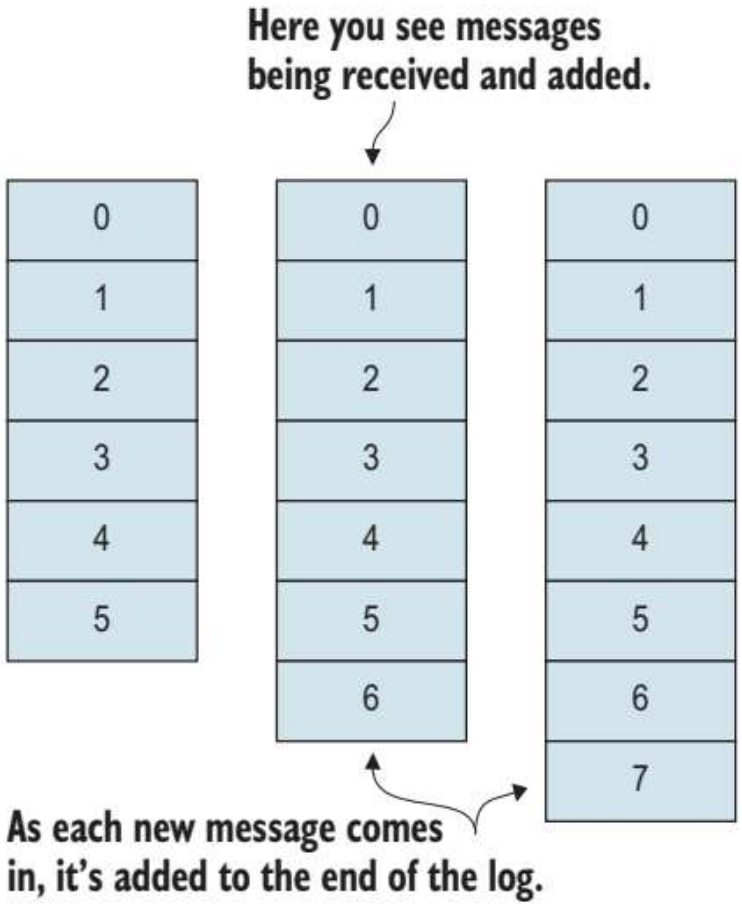
Here you see messages being received and added.
0
1
2
3
4
5
0
1
2
3
4
5
6
0
1
2
3
4
5
6
7
As each new message comes in, it's added to the end of the log.

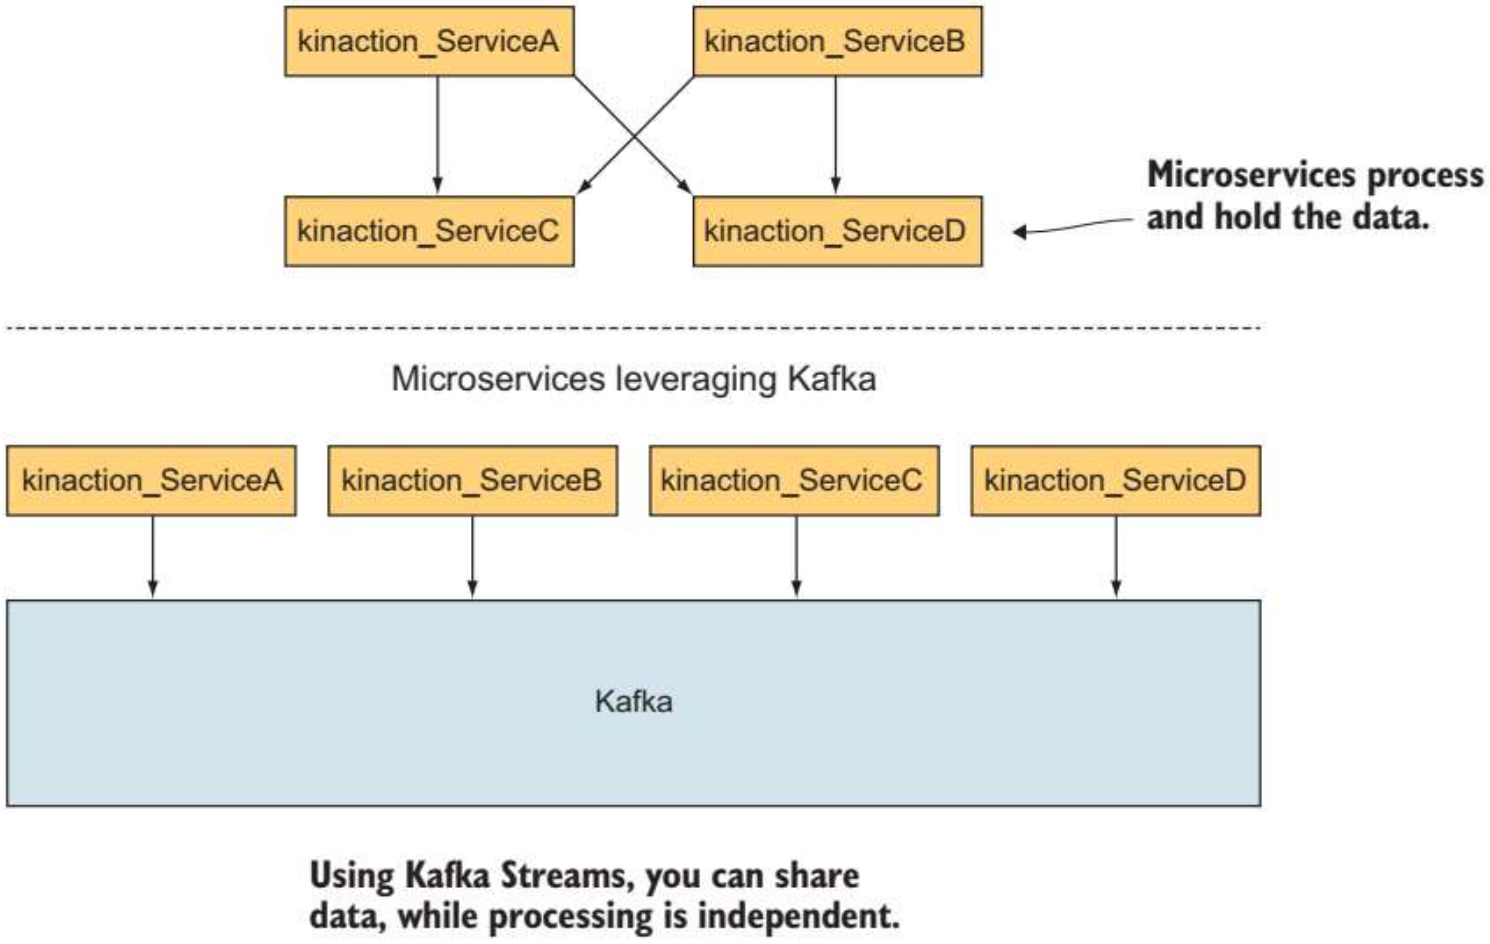
kinaction_ServiceA
kinaction_ServiceB
kinaction_ServiceC
kinaction_ServiceD
Microservices process and hold the data.
Microservices leveraging Kafka
kinaction_ServiceA
kinaction_ServiceB
kinaction_ServiceC
kinaction_ServiceD
Kafka
Using Kafka Streams, you can share data, while processing is independent.

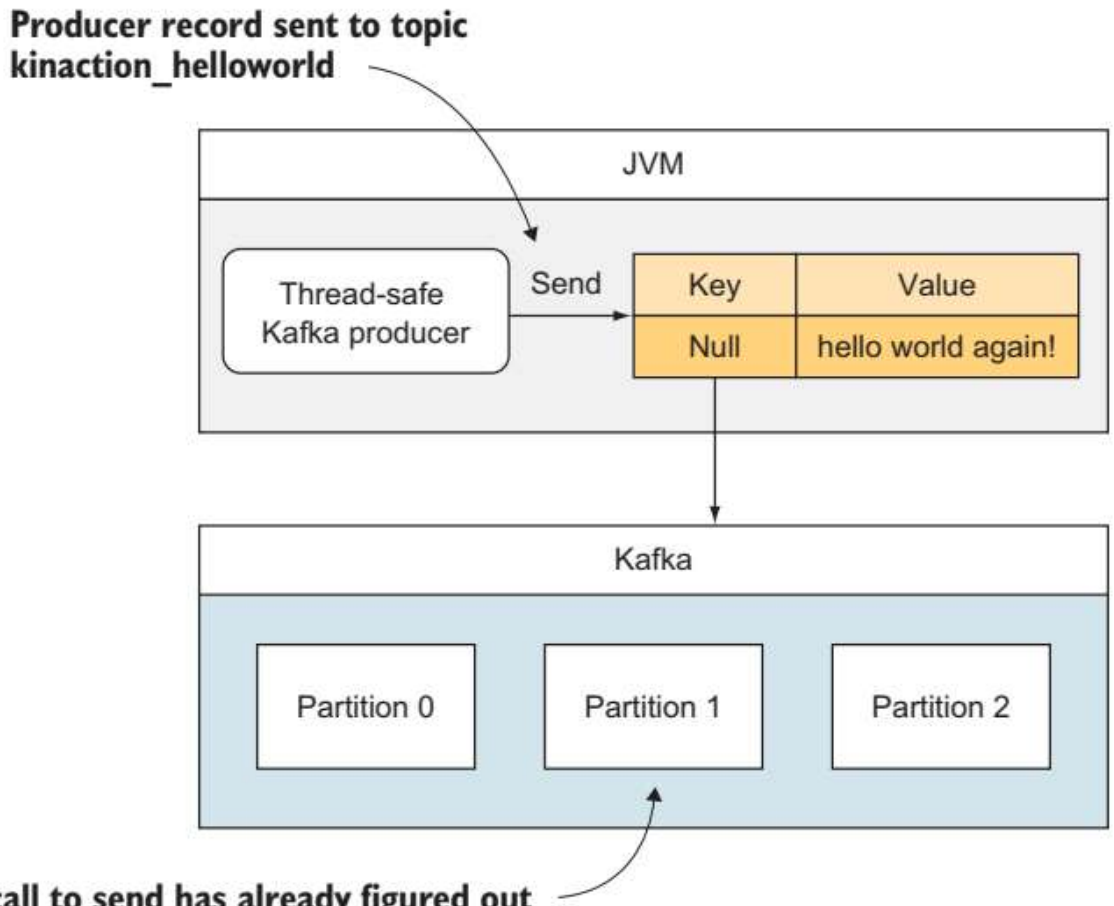

Producer flow

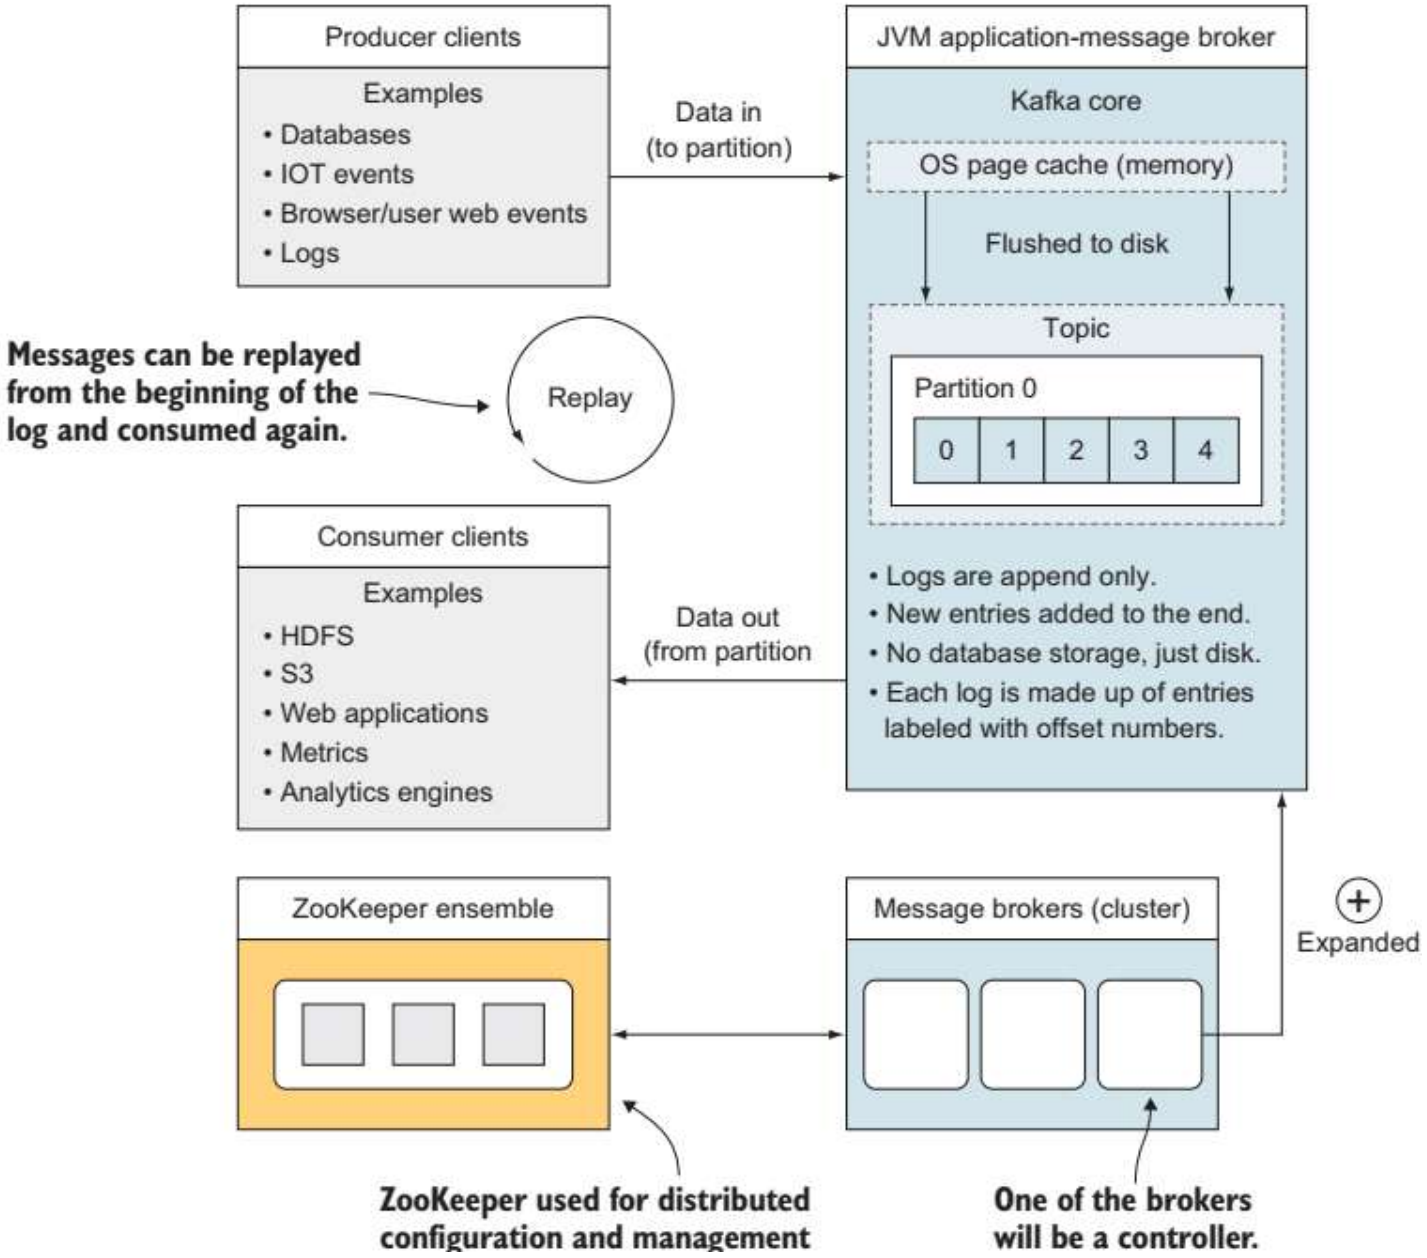
Producer clients
Examples
• Databases
• IOT events
• Browser/user web events
• Logs
Data in
(to partition)
JVM application-message broker
Kafka core
OS page cache (memory)
Flushed to disk
Topic
Partition 0
0
1
2
3
4
• Logs are append only.
• New entries added to the end.
• No database storage, just disk.
• Each log is made up of entries labeled with offset numbers.
Messages can be replayed from the beginning of the log and consumed again.
Replay
Consumer clients
Examples
• HDFS
• S3
• Web applications
• Metrics
• Analytics engines
Data out
(from partition
ZooKeeper ensemble
Message brokers (cluster)
Expanded
ZooKeeper used for distributed configuration and management
One of the brokers will be a controller.

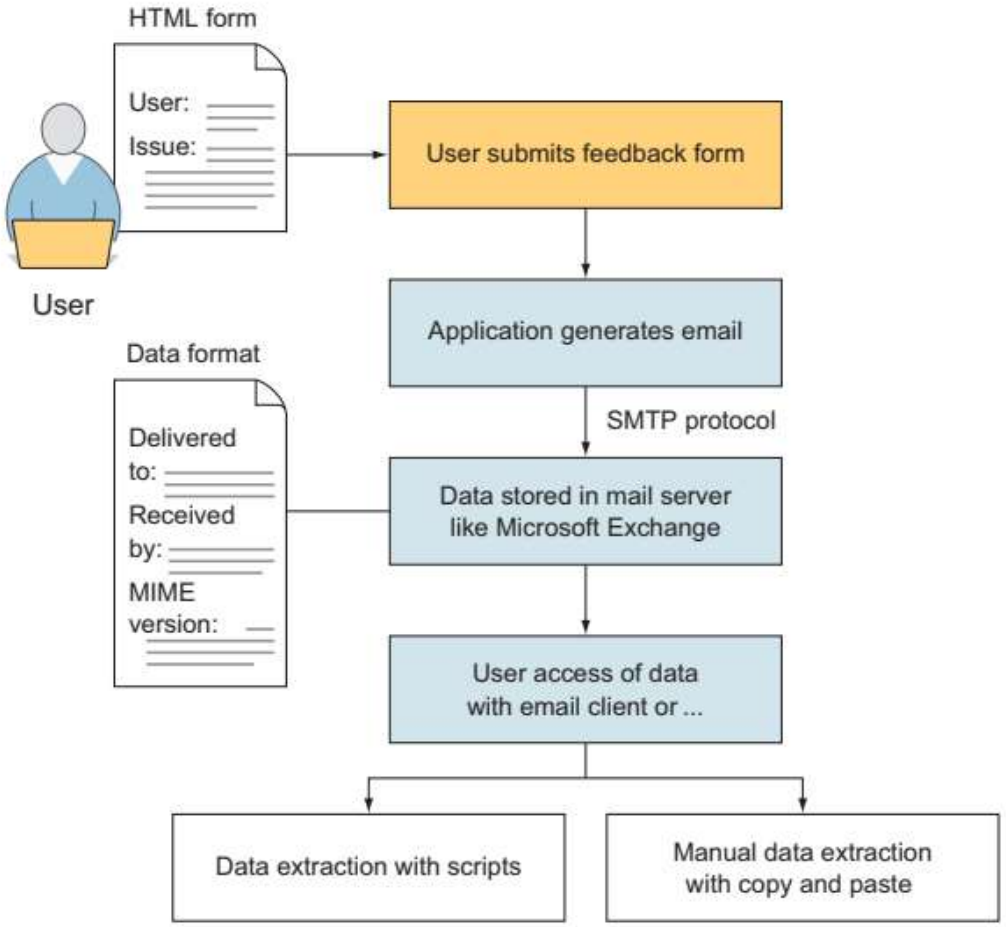

Sending data in email

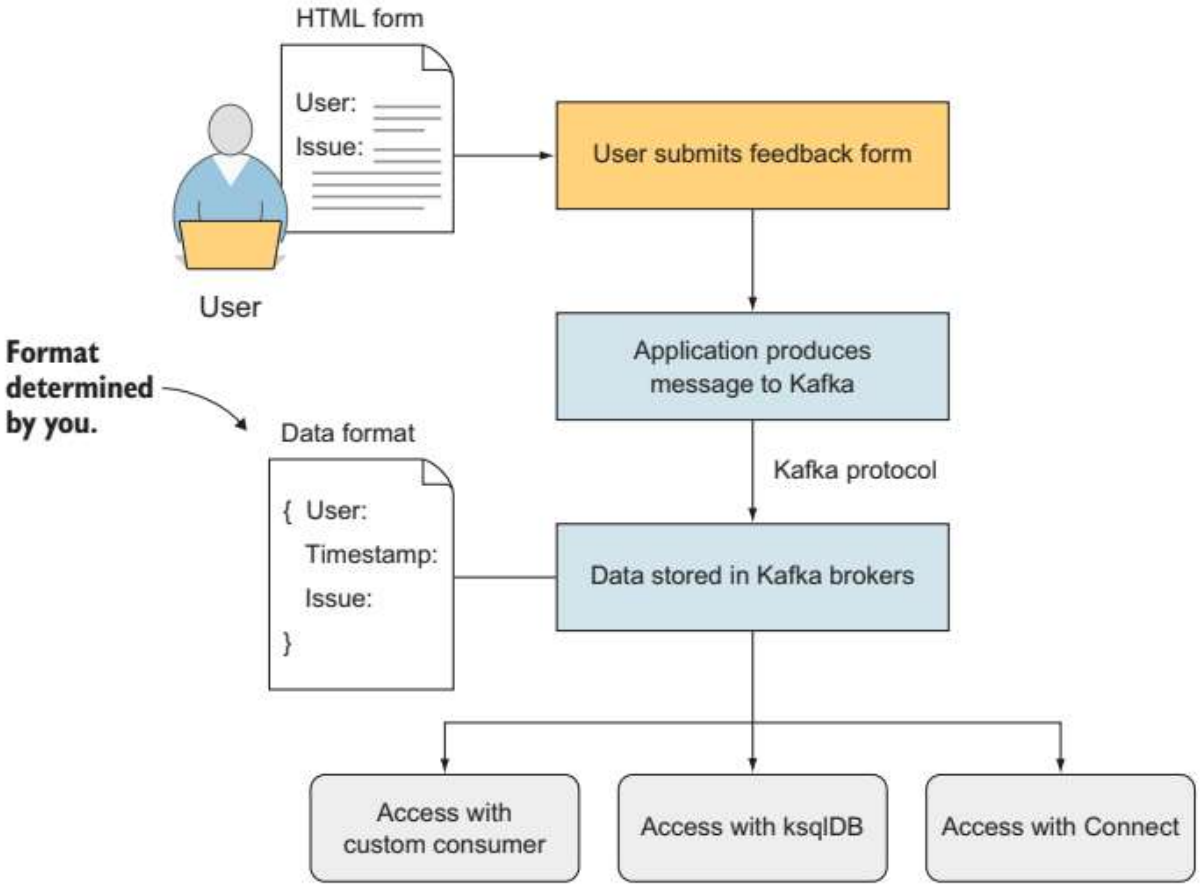

Sending data to Kafka

Important producer configurations

Key	Purpose
`acks`	Number of replica acknowledgments that a producer requires before success is established
`bootstrap.servers`	One or more Kafka brokers to connect for startup
`value.serializer`	The class that's used for serialization of the value
`key.serializer`	The class that's used for serialization of the key

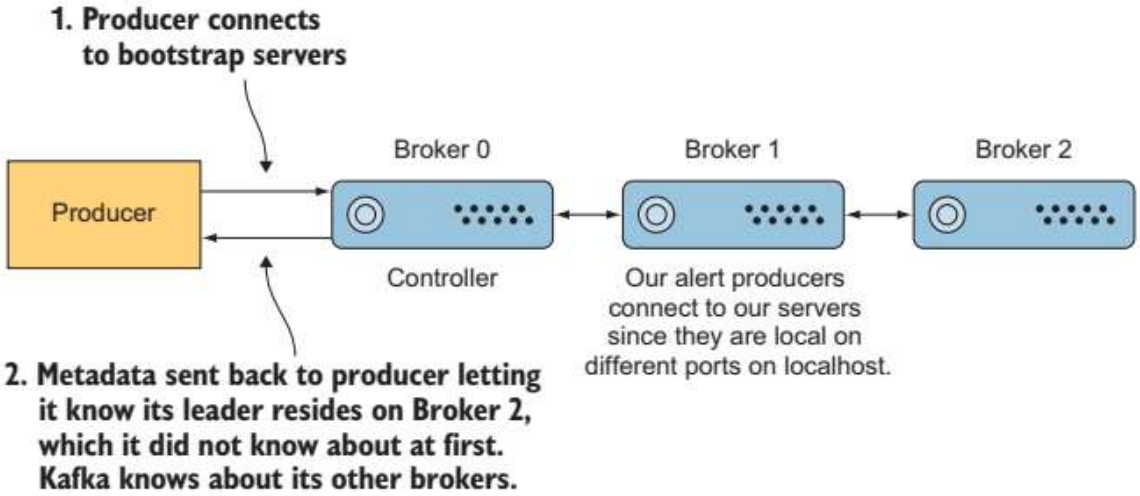
1. Producer connects to bootstrap servers
Broker 0
Broker 1
Broker 2
Producer
Controller
Our alert producers connect to our servers since they are local on different ports on localhost.
2. Metadata sent back to producer letting it know its leader resides on Broker 2, which it did not know about at first. Kafka knows about its other brokers.
Bootstrap servers

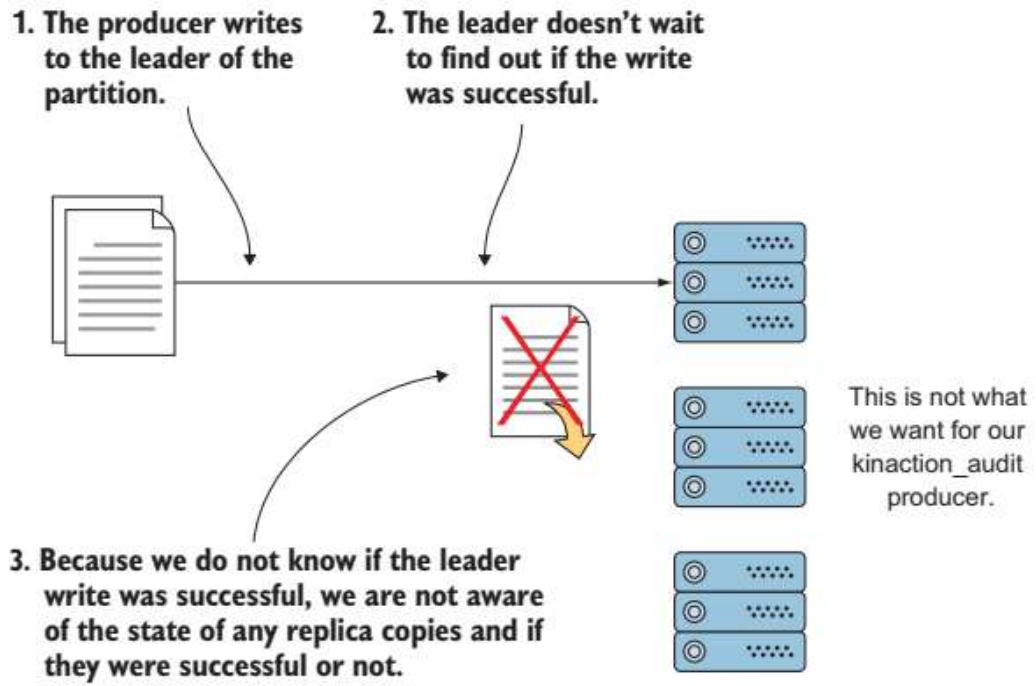

The property `acks` equals `0`.

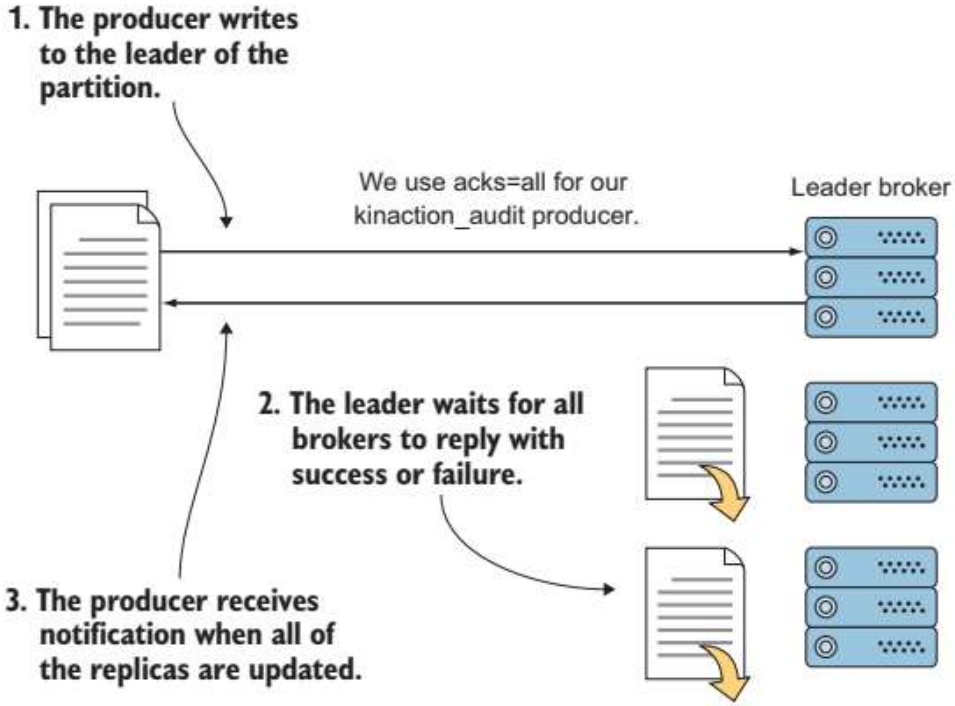

The property `acks` equals `all`.

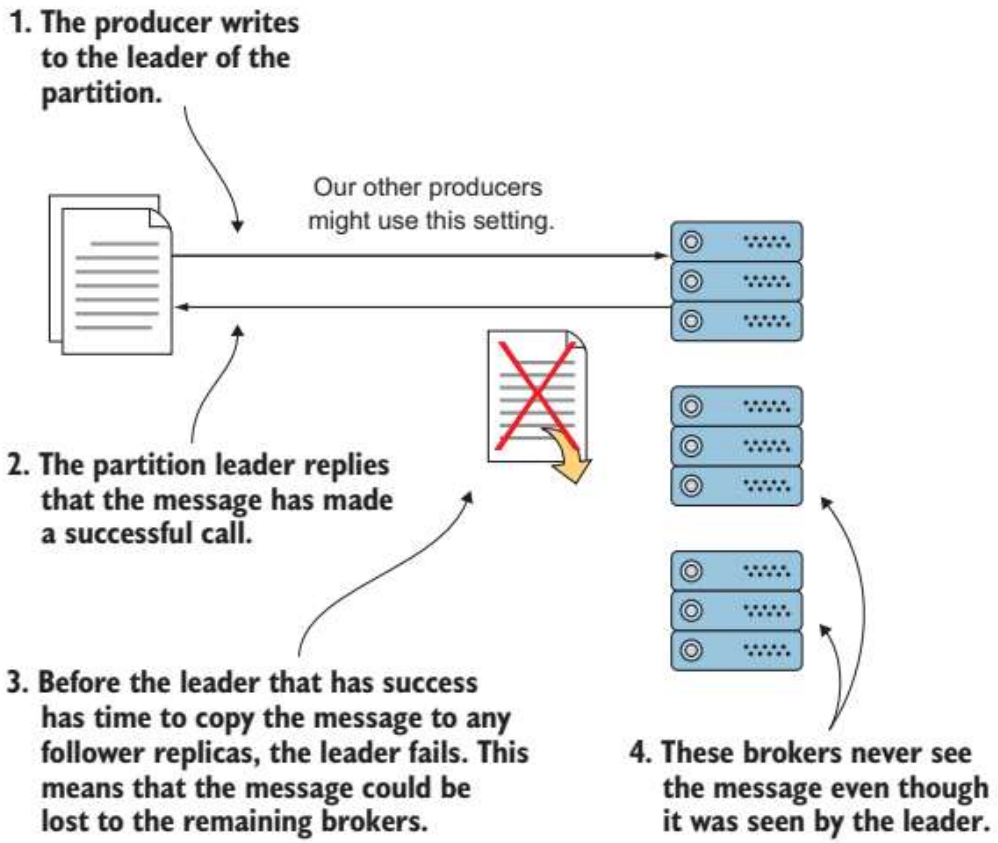

The property `acks` equals `1`.

Consumer configuration

Key	Purpose
`bootstrap.servers`	One or more Kafka brokers to connect on startup
`value.deserializer`	Needed for deserialization of the value
`key.deserializer`	Needed for deserialization of the key
`group.id`	A name that's used to join a consumer group
`client.id`	An ID to identify a user
`heartbeat.interval.ms`	Interval for consumer's pings to the group coordinator

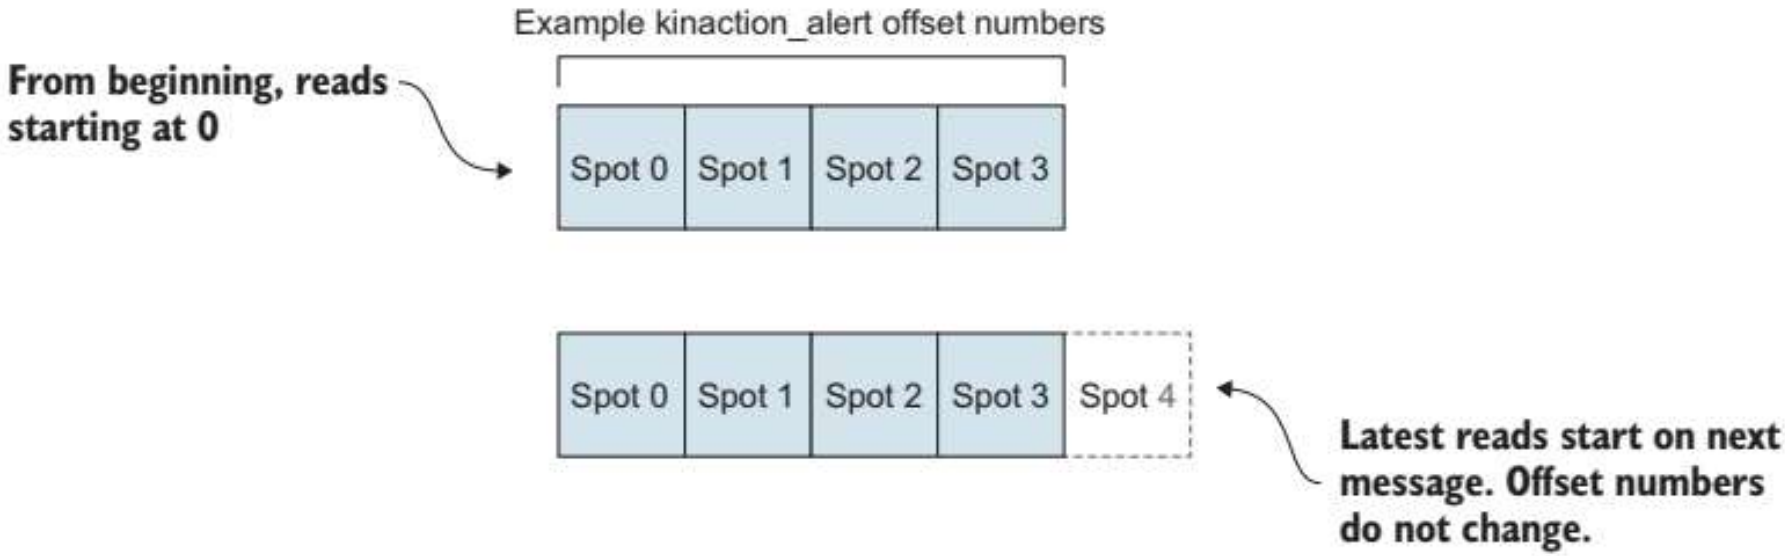

Kafka offsets [6]

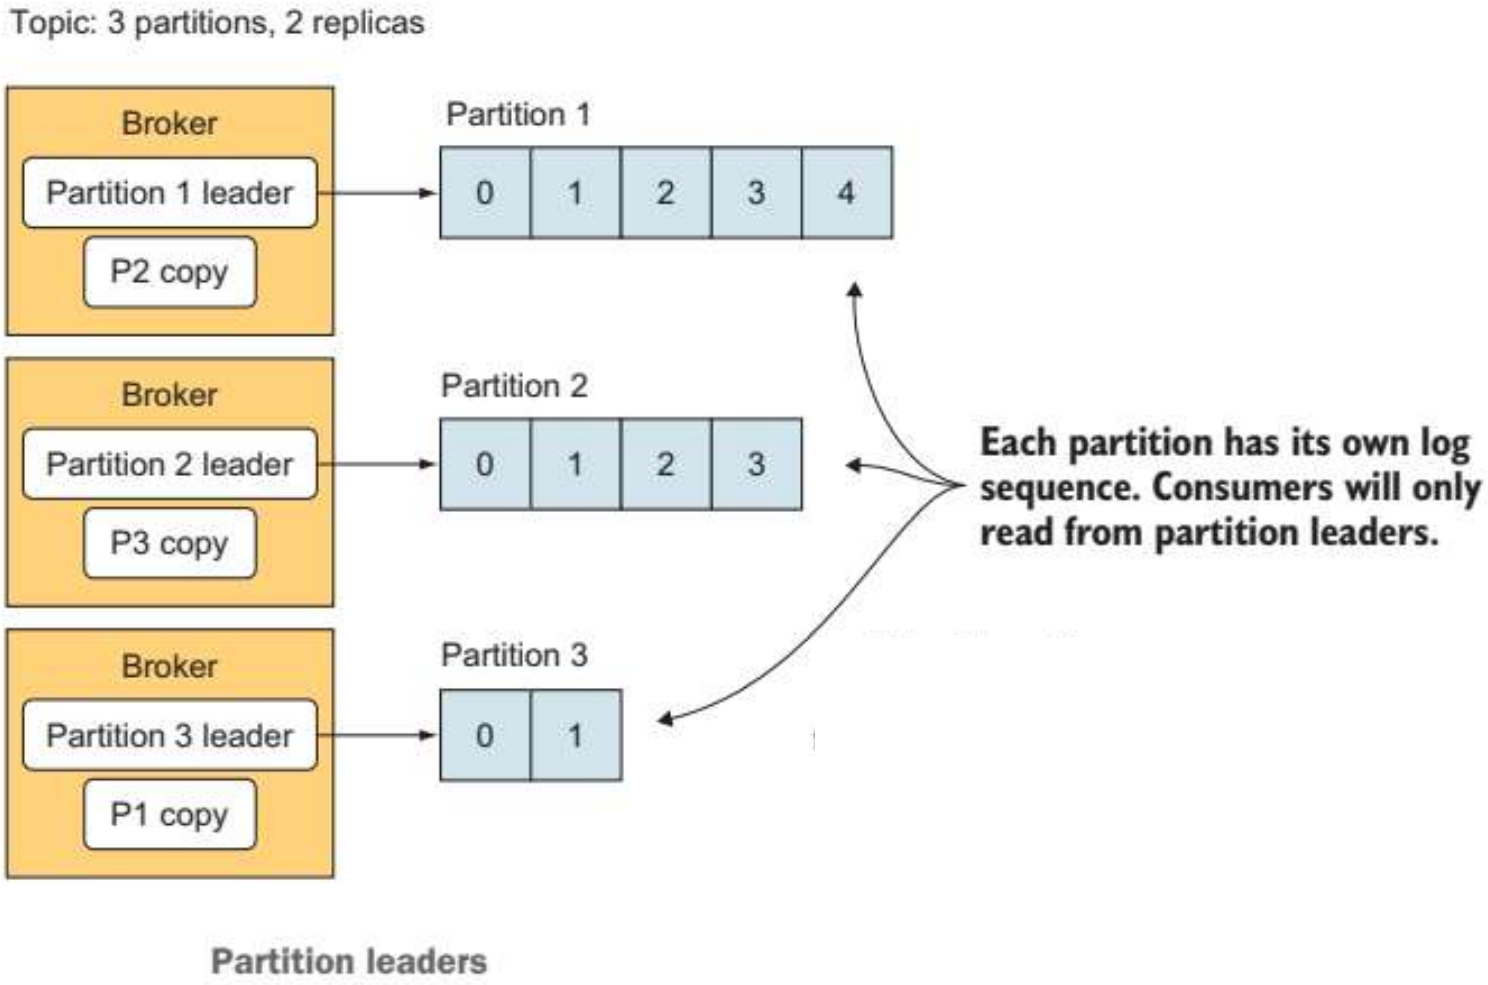
Topic: 3 partitions, 2 replicas
Broker
Partition 1 leader
P2 copy
Partition 1
0
1
2
3
4
Broker
Partition 2 leader
P3 copy
Partition 2
0
1
2
3
Broker
Partition 3 leader
P1 copy
Partition 3
0
1
Each partition has its own log sequence. Consumers will only read from partition leaders.
Partition leaders

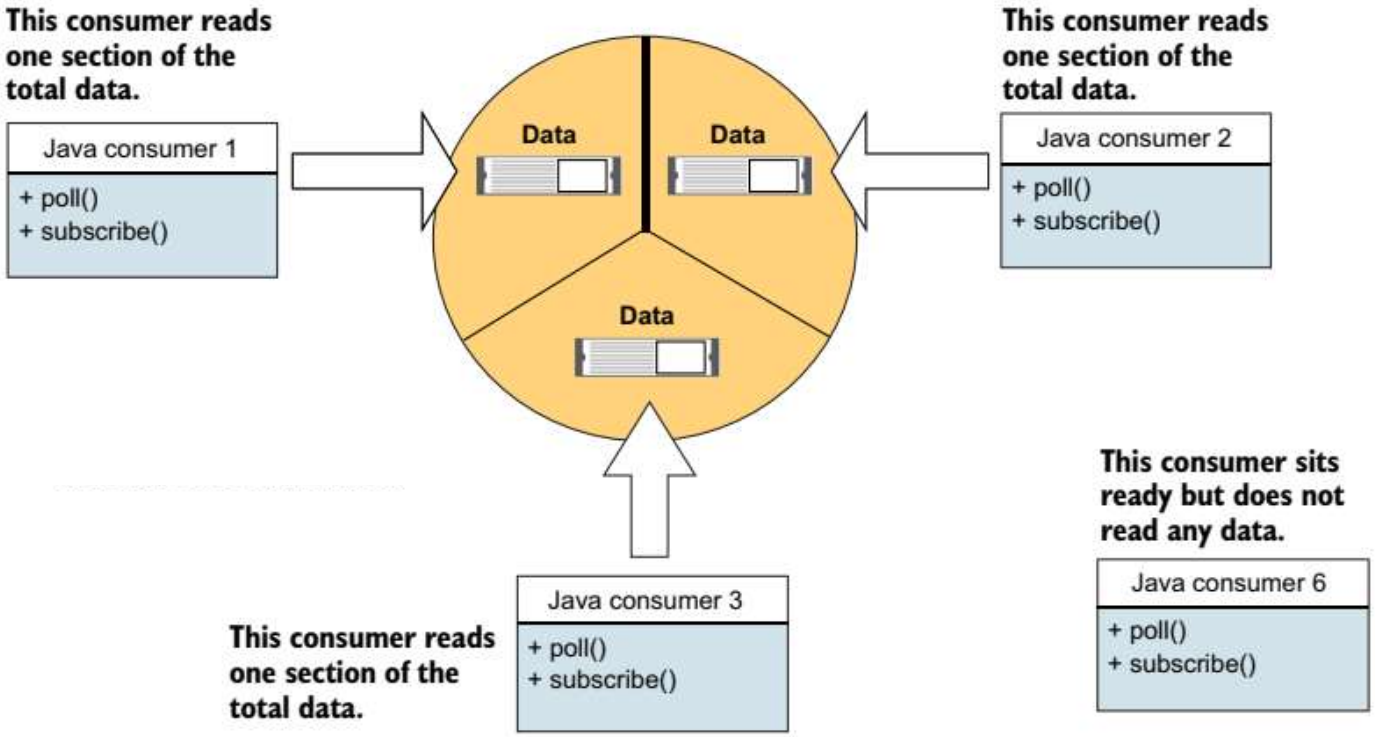

An extra Kafka consumer

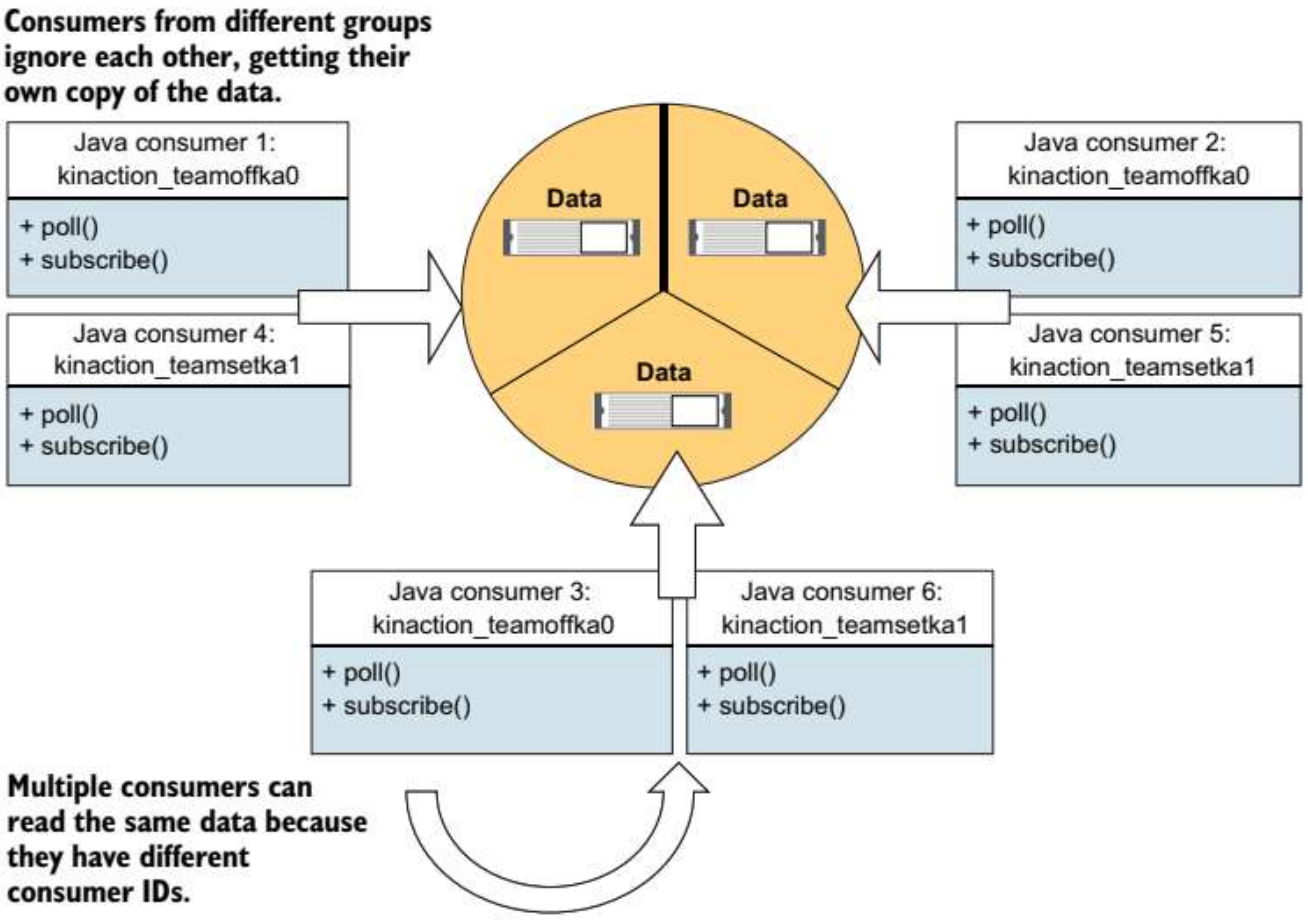

Consumers in separate groups [12]

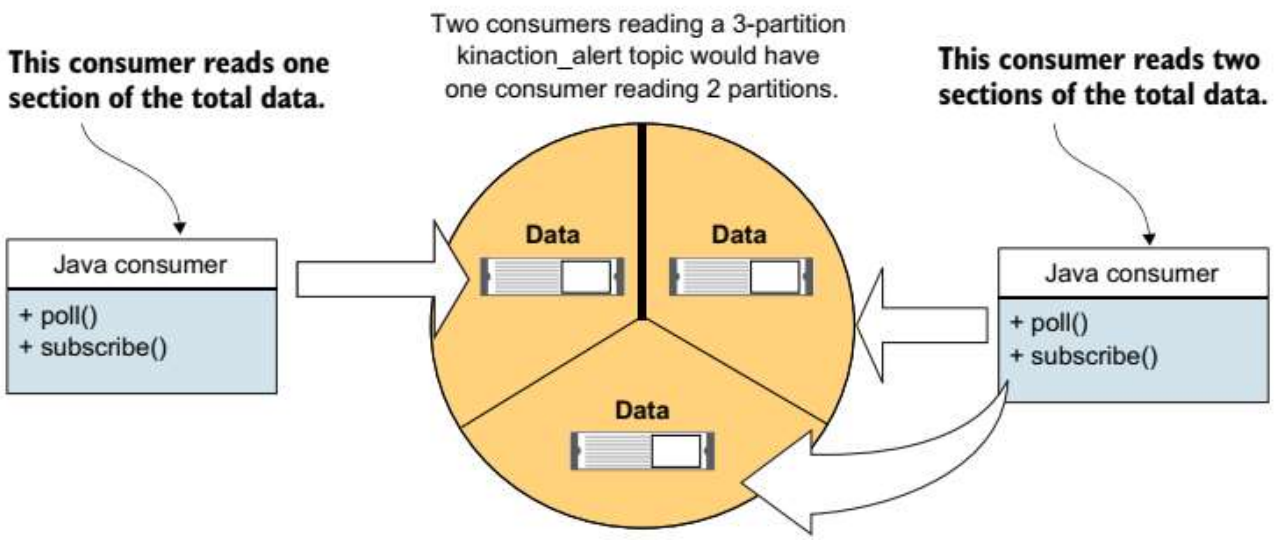

Kafka consumers in a group

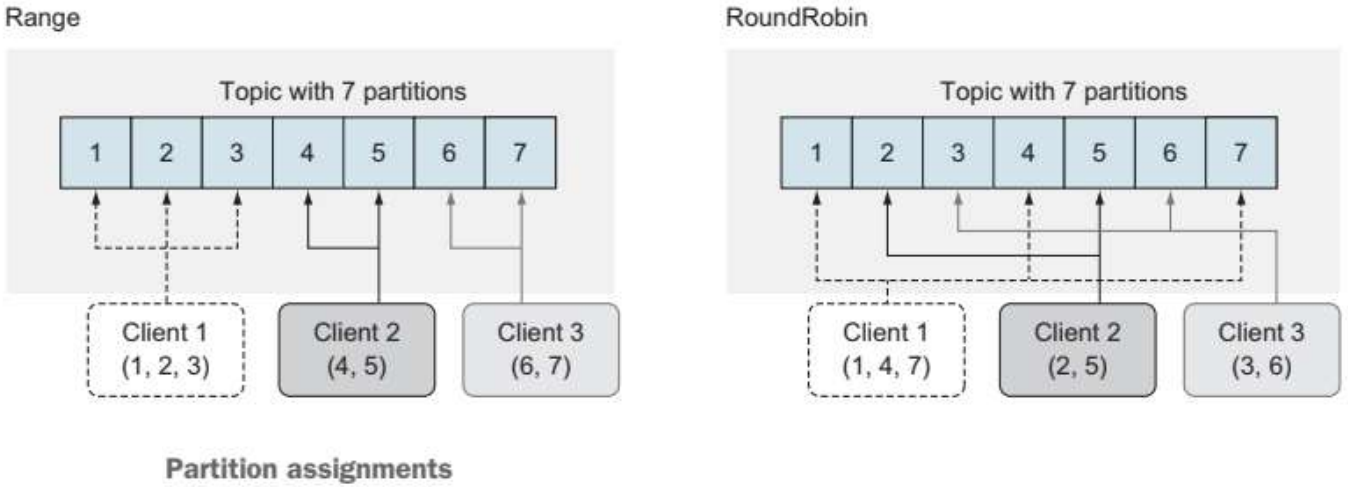

Partition assignments

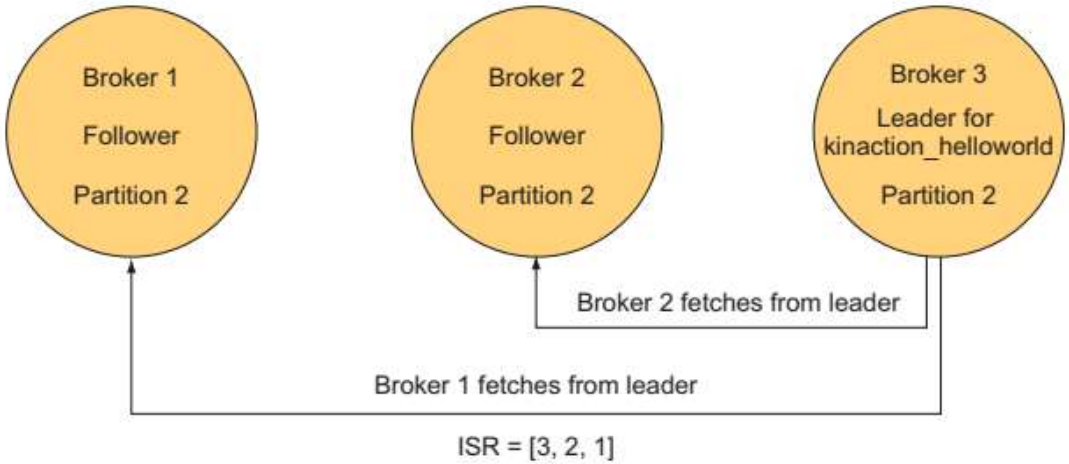
Broker 1
Follower
Partition 2
Broker 2
Follower
Partition 2
Broker 3
Leader for
kinaction_helloworld
Partition 2
Broker 2 fetches from leader
Broker 1 fetches from leader
ISR = [3, 2, 1]

Leader

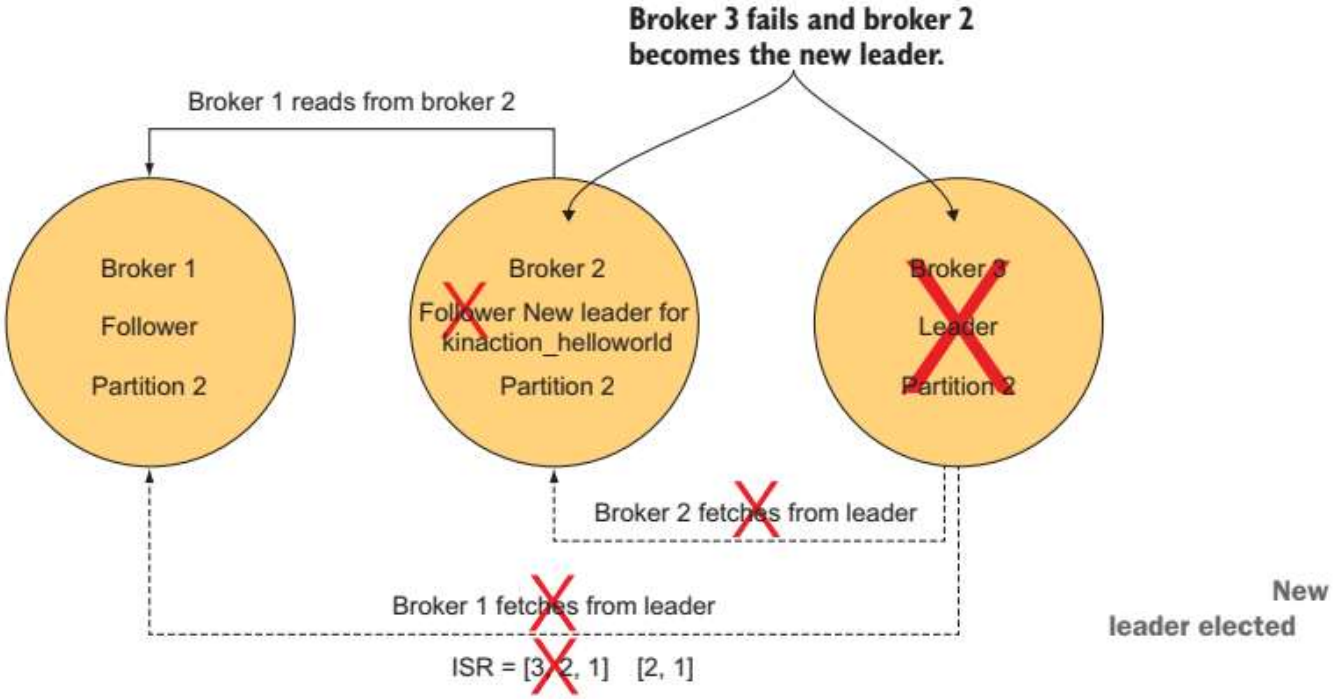
Broker 3 fails and broker 2 becomes the new leader.
Broker 1 reads from broker 2
Broker 1
Follower
Partition 2
Broker 2
Follower New leader for kinaction_helloworld
Partition 2
Broker 3
Leader
Partition 2
Broker 2 fetches from leader
Broker 1 fetches from leader
ISR = [3, 2, 1] [2, 1]
New leader elected

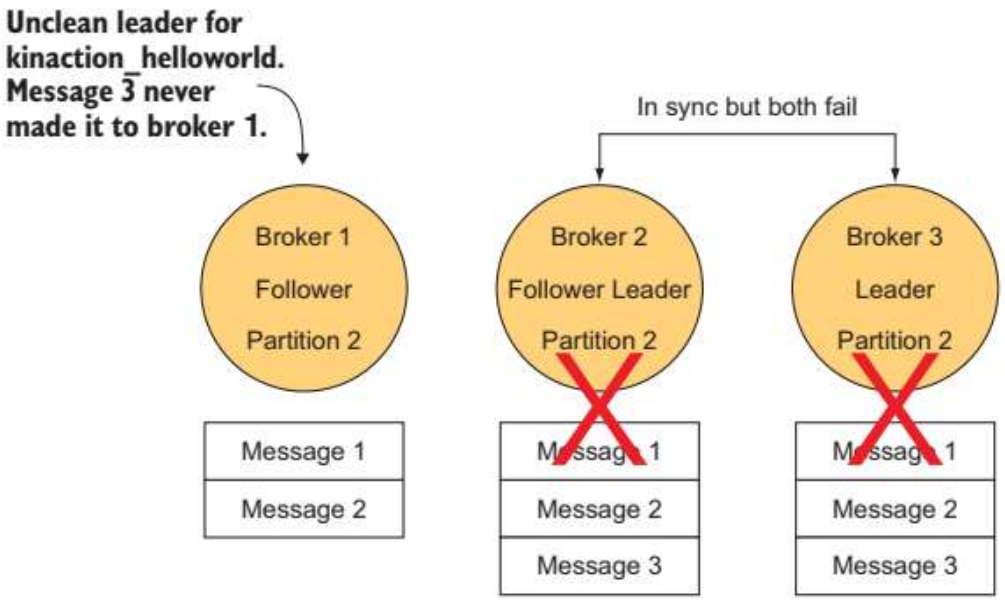

Unclean leader election

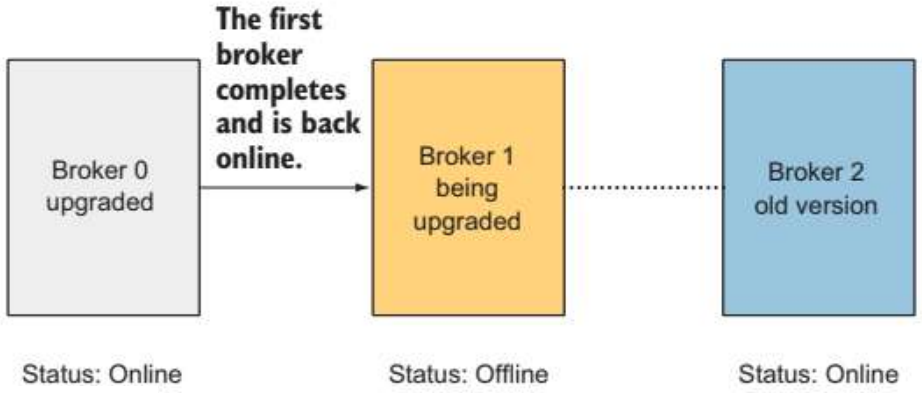

Rolling restart

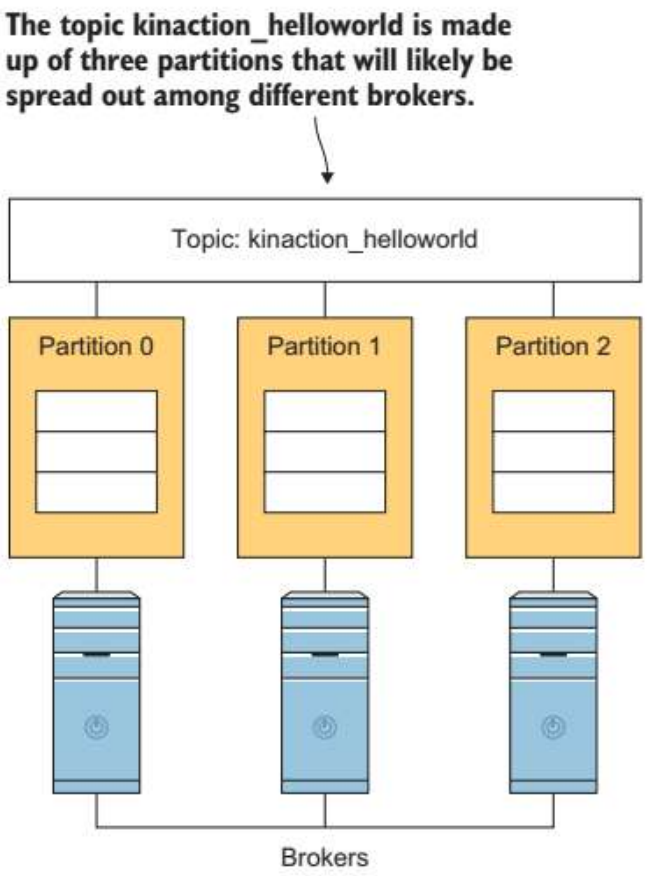

Example topic with partitions

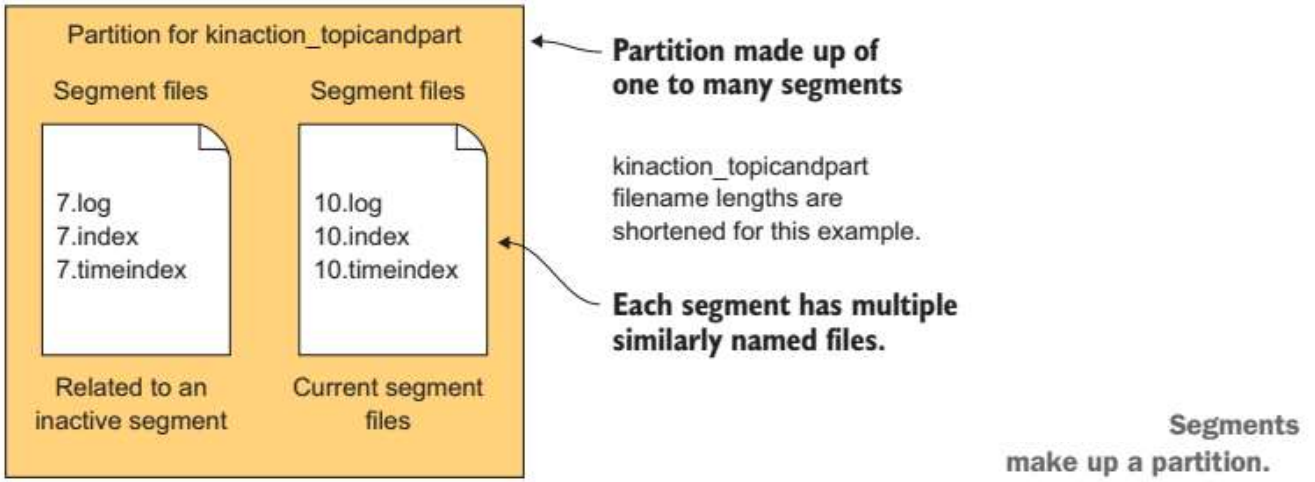

Segments make up a partition.

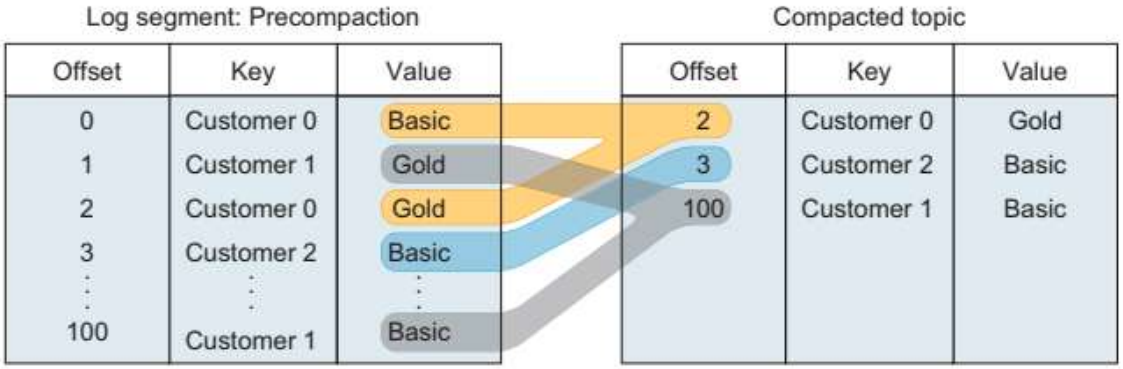

Compaction in general

Broker retention configuration

Key	Purpose
`log.retention.bytes`	The largest size threshold in bytes for deleting a log.
`log.retention.ms`	The length in milliseconds a log will be maintained before being deleted.
`log.retention.minutes`	Length before deletion in minutes. `log.retention.ms` is used as well if both are set.
`log.retention.hours`	Length before deletion in hours. `log.retention.ms` and `log.retention.minutes` would be used before this value if either of those are set.

Scenario 1: no advertised listeners. Producer client starts and requests metadata from bootstrap server.

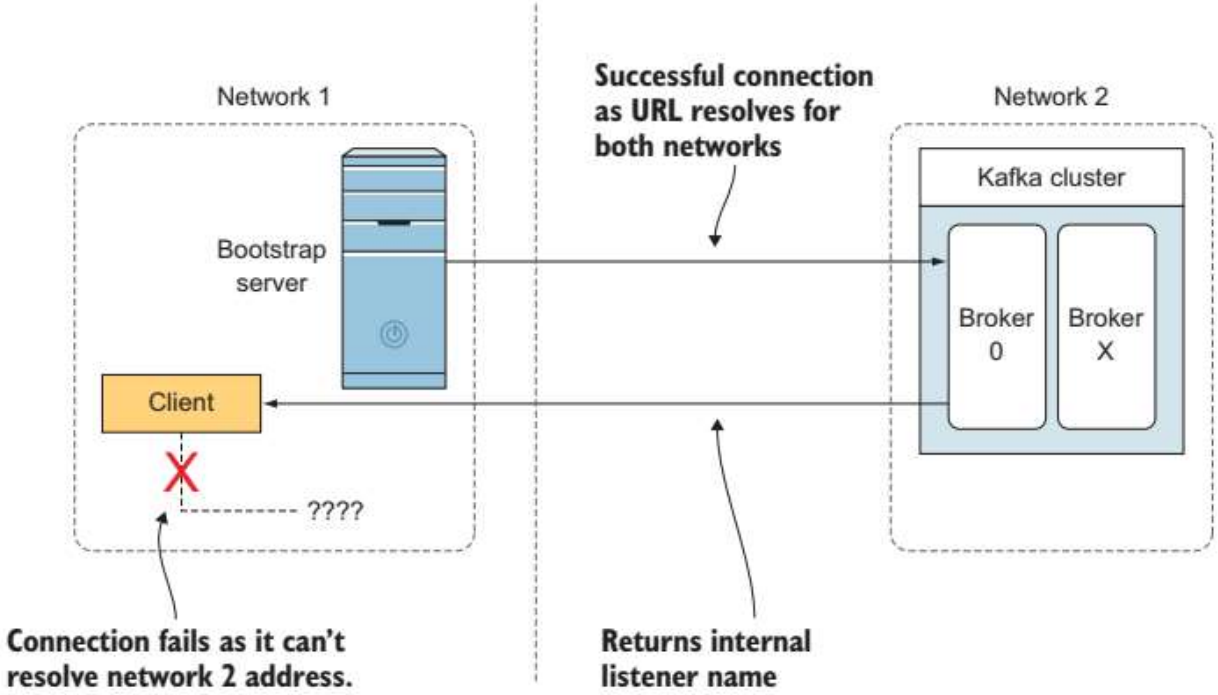

Scenario 2: advertised listeners with URL resolved by both networks. Producer client requests metadata.

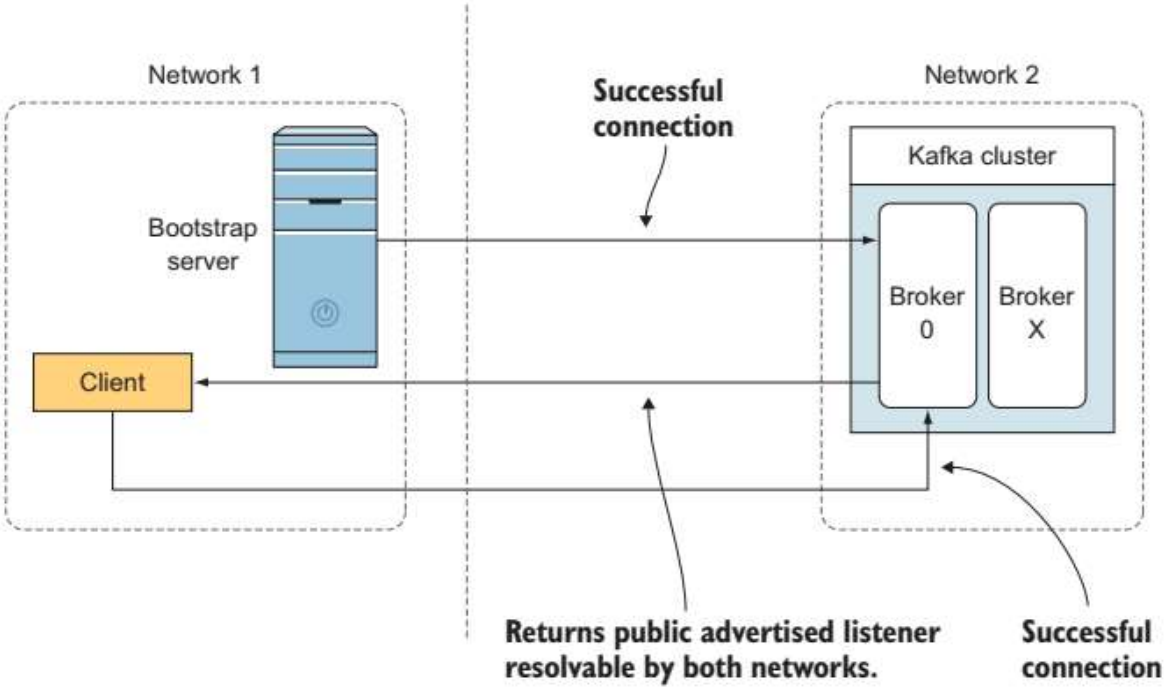

Thanks